

# Model

July 6, 2025

## 1 Modeling

### 1.1 Setup Code

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV, train_test_split, \
    cross_val_score, TimeSeriesSplit
from sklearn.linear_model import ElasticNet, Lasso, Ridge
from sklearn.ensemble import RandomForestRegressor
from sklearn.decomposition import PCA
from xgboost import XGBRegressor
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import statsmodels.api as sm

pd.set_option("display.max_rows", None)      # Show all rows
pd.set_option("display.max_columns", None)   # Show all columns
pd.set_option("display.width", 2000)         # Prevent line wrapping
pd.set_option("display.max_colwidth", None) # Don't truncate cell conte

[2]: data = pd.read_pickle('C:/Code/Git Repositories/Football/Football/rawdata_clean.
    pickle')

[3]: df = data.select_dtypes(include=['int64', 'float64'])
df = df.dropna()
sorted_df = df.sort_values(by=['season', 'week'])
X = sorted_df.drop(columns=['margin', 'away_score', 'home_score'])
y = sorted_df['margin']

split_idx = int(len(X) * 0.8)
tss_X_train, tss_X_test = X.iloc[:split_idx], X.iloc[split_idx:]
tss_y_train, tss_y_test = y.iloc[:split_idx], y.iloc[split_idx:]
```

```
X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=.2,
    ↪random_state=369, shuffle=False)
X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp, test_size=.
    ↪.25, random_state=369, shuffle=False)
```

## 1.2 Model Selection

Lasso Regression is continuously the best performer

```
[4]: elastic_pipe = Pipeline(steps=[('scaler', StandardScaler()), ('ElasticNet',
    ↪ElasticNet())])
ridge_pipe = Pipeline(steps=[('scaler', StandardScaler()), ('Ridge', Ridge())])
lasso_pipe = Pipeline(steps=[('scaler', StandardScaler()), ('Lasso', Lasso())])
xgb = XGBRegressor()
rand_forest = RandomForestRegressor()
tscv = TimeSeriesSplit(n_splits=5)

elastic_cv = cross_val_score(elastic_pipe, X_train, y_train, cv=tscv,
    ↪scoring='r2')
ridge_cv = cross_val_score(ridge_pipe, X_train, y_train, cv=tscv, scoring='r2')
lasso_cv = cross_val_score(lasso_pipe, X_train, y_train, cv=tscv, scoring='r2')
xgb_cv = cross_val_score(xgb, X_train, y_train, cv=tscv, scoring='r2')
rand_forest_cv = cross_val_score(rand_forest, X_train, y_train, cv=tscv,
    ↪scoring='r2')

print(f'elastic_cv score average: {elastic_cv.mean() * 100: .2f}. Standard
    ↪deviation is {elastic_cv.std() * 100: .2f}')
print(f'ridge_cv score average: {ridge_cv.mean() * 100: .2f}. Standard
    ↪deviation is {ridge_cv.std() * 100: .2f}')
print(f'lasso_cv score average: {lasso_cv.mean() * 100: .2f}. Standard
    ↪deviation is {lasso_cv.std() * 100: .2f}')
print(f'xgb_cv score average: {xgb_cv.mean() * 100: .2f}. Standard deviation is
    ↪{xgb_cv.std() * 100: .2f}')
print(f'rand_forest_cv score average: {rand_forest_cv.mean() * 100: .2f}.
    ↪Standard deviation is {rand_forest_cv.std() * 100: .2f}')
```

```
elastic_cv score average: 74.16. Standard deviation is 1.00
ridge_cv score average: 73.09. Standard deviation is 8.21
lasso_cv score average: 75.81. Standard deviation is 1.29
xgb_cv score average: 72.92. Standard deviation is 3.56
rand_forest_cv score average: 72.57. Standard deviation is 2.93
```

## 1.3 Model Tuning

~5.5% increase from initial training

```
[5]: lasso_param_grid = {
    'Lasso__alpha': [0.01, 0.1, 0.3, 0.5, 0.7, 1],
```

```

        'Lasso__fit_intercept': [True, False],
        'Lasso__max_iter': [1000, 3000, 10000]
    }

    search = GridSearchCV(lasso_pipe, param_grid=lasso_param_grid, scoring='r2').
        ↪fit(X_train, y_train)

    print(f'Best parameters: {search.best_params_}')
    print(f'Best score: {search.best_score_ * 100: .2f}')
```

```

Best parameters: {'Lasso__alpha': 0.1, 'Lasso__fit_intercept': True,
'Lasso__max_iter': 1000}
Best score: 80.07
```

```

[6]: params = search.best_params_
    clean_params = {k.replace('Lasso__', ''): v for k, v in params.items()}
    tuned_lasso = Lasso(**clean_params).fit(X_train, y_train)

    print(f'Tuned Lasso score: {tuned_lasso.score(X_train, y_train) * 100: .2f}')
```

```
Tuned Lasso score: 81.24
```

## 1.4 Model Validation & Testing

Model generalizes very well

### 1.4.1 Validation

```

[7]: predictions = tuned_lasso.predict(X_val)
    mae = mean_absolute_error(predictions, y_val)
    rmse = np.sqrt(mean_squared_error(predictions, y_val))
    score = r2_score(predictions, y_val)

    print(f'mae is {mae: .2f}')
    print(f'rmse is {rmse: .2f}')
    print(f'r2 score is {score * 100: .2f}')
```

```

mae is 4.99
rmse is 6.29
r2 score is 76.90
```

### 1.4.2 Test

```

[8]: predictions = tuned_lasso.predict(X_test)
    mae = mean_absolute_error(predictions, y_test)
    rmse = np.sqrt(mean_squared_error(predictions, y_test))
    score = r2_score(predictions, y_test)

    print(f'mae is {mae: .2f}')
```


```
print(f'rmse is {rmse: .2f}')
print(f'r2 score is {score * 100: .2f}')
```

```
mae is 4.99
rmse is 6.43
r2 score is 75.05
```

## 1.5 Residual Analysis


```
[9]: predictions = tuned_lasso.predict(X_test)
     residuals = y_test - predictions
```

```
[10]: plt.figure(figsize=(6,4))
      plt.scatter(predictions, residuals, alpha=.5)
      plt.axhline(0, color='red', linestyle='--')
      plt.ylabel('Residuals')
      plt.xlabel('Fitted Values')
      plt.title('Residuals vs. Fitted')
      plt.show()
```



```
[11]: sns.histplot(residuals, kde=True)
      plt.xlabel('Residuals')
      plt.title('Normality Check. Histogram of Residuals')
```

```
plt.show()
```



```
[12]: dw = sm.stats.durbin_watson(residuals)

print(f'Durbin-Watson score: {dw: .2f}')
```

Durbin-Watson score: 1.96

```
[13]: top_10_thresh = sorted_df['margin'].quantile(0.90)
bottom_10_thresh = sorted_df['margin'].quantile(0.10)

top_10_df = sorted_df[sorted_df['margin'] >= top_10_thresh]
bottom_10_df = sorted_df[sorted_df['margin'] <= bottom_10_thresh]

features_df = sorted_df.drop(columns=['margin', 'home_score', 'away_score'])
features = features_df.columns

top_10_preds = tuned_lasso.predict(top_10_df[features])
bottom_10_preds = tuned_lasso.predict(bottom_10_df[features])
```

```

print('Top 10%:')
print('RMSE:', mean_squared_error(top_10_df['margin'], top_10_preds,
    ↪squared=False))
print('MAE:', mean_absolute_error(top_10_df['margin'], top_10_preds))

print('\nBottom 10%:')
print('RMSE:', mean_squared_error(bottom_10_df['margin'], bottom_10_preds,
    ↪squared=False))
print('MAE:', mean_absolute_error(bottom_10_df['margin'], bottom_10_preds))

```


Top 10%:  
 RMSE: 8.901360406525646  
 MAE: 7.007829147824513

Bottom 10%:  
 RMSE: 8.249468437213423  
 MAE: 6.635713699781495


```

[14]: plt.scatter(top_10_df['margin'], top_10_preds, alpha=0.6)
plt.plot([-50, 50], [-50, 50], color='red', linestyle='--')
plt.xlabel('Actual Margin (Top 10%)')
plt.ylabel('Predicted Margin')
plt.title('Prediction Accuracy on Top 10% Margin Games')
plt.show()

```



```
[15]: plt.scatter(bottom_10_df['margin'], bottom_10_preds, alpha=0.6)
plt.plot([-50, 50], [-50, 50], color='red', linestyle='--')
plt.xlabel('Actual Margin (Top 10%)')
plt.ylabel('Predicted Margin')
plt.title('Prediction Accuracy on Top 10% Margin Games')
plt.show()
```



```
[ ]: window_size = 816 # adjust to your dataset size
test_size = 272
errors = {}

for start in range(0, len(sorted_df) - window_size - test_size):
    train = sorted_df.iloc[start : start + window_size]
    test = sorted_df.iloc[start + window_size : start + window_size + test_size]

    X_train = train[features]
    y_train = train['margin']
    X_test = test[features]
    y_test = test['margin']

    model = Lasso(**clean_params)
    model.fit(X_train, y_train)
    preds = model.predict(X_test)

    mae = mean_absolute_error(y_test, preds)
    r2 = r2_score(y_test, preds)
    rmse = np.sqrt(mean_squared_error(y_test, preds))
```



```

errors[start] = {'mae' : mae, 'rmse' : rmse, 'r2' : r2}

mae_mean = np.mean([metrics['mae'] for metrics in errors.values()])
rmse_mean = np.mean([metrics['rmse'] for metrics in errors.values()])
r2_mean = np.mean([metrics['r2'] for metrics in errors.values()])

print(f"Past three seasons Mean MAE: {mae_mean:.2f}")
print(f"Past three seasons Mean RMSE: {rmse_mean:.2f}")
print(f"Past three seasons Mean R2: {r2_mean * 100 :.2f}%")

```

Past three seasons Mean MAE: 5.39  
 Past three seasons Mean RMSE: 6.88  
 Past three seasons Mean R2: 77.60%

```

[17]: window_size = 6300 # adjust to your dataset size
      test_size = 272
      errors = {}

      for start in range(0, len(sorted_df) - window_size - test_size):
          train = sorted_df.iloc[start : start + window_size]
          test = sorted_df.iloc[start + window_size : start + window_size + test_size]

          X_train = train[features]
          y_train = train['margin']
          X_test = test[features]
          y_test = test['margin']

          model = Lasso(**clean_params)
          model.fit(X_train, y_train)
          preds = model.predict(X_test)

          mae = mean_absolute_error(y_test, preds)
          r2 = r2_score(y_test, preds)
          rmse = np.sqrt(mean_squared_error(y_test, preds))
          errors[start] = {'mae' : mae, 'rmse' : rmse, 'r2' : r2}

      mae_mean = np.mean([metrics['mae'] for metrics in errors.values()])
      rmse_mean = np.mean([metrics['rmse'] for metrics in errors.values()])
      r2_mean = np.mean([metrics['r2'] for metrics in errors.values()])

      print(f"All game data Mean MAE: {mae_mean:.2f}")
      print(f"All game data Mean RMSE: {rmse_mean:.2f}")
      print(f"All game data Mean R2: {r2_mean * 100 :.2f}%")

```

All game data Mean MAE: 4.82  
 All game data Mean RMSE: 6.21  
 All game data Mean R2: 81.07%